

meteoregresion

October 16, 2020

Predicción humedad a las 3 pm a partir de los features calculados entre las 8:55 am y 9:04 am por método de regresión

```
[3]: import pandas as pd
from sklearn import preprocessing
from sklearn.metrics import accuracy_score
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LinearRegression
from math import sqrt

import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

Cargamos los datos

```
[4]: data = pd.read_csv('meteo/diario.csv')
data.columns
data.shape
```

```
[4]: (1095, 11)
```

```
[7]: data.head(10)
```

```
[7]:
```

	number	air_pressure_9am	air_temp_9am	avg_wind_direction_9am	\
0	0	918.060000	74.822000	271.100000	
1	1	917.347688	71.403843	101.935179	
2	2	923.040000	60.638000	51.000000	
3	3	920.502751	70.138895	198.832133	
4	4	921.160000	44.294000	277.800000	
5	5	915.300000	78.404000	182.800000	
6	6	915.598868	70.043304	177.875407	
7	7	918.070000	51.710000	242.400000	
8	8	920.080000	80.582000	40.700000	
9	9	915.010000	47.498000	163.100000	

```
avg_wind_speed_9am max_wind_direction_9am max_wind_speed_9am \
```

0	2.080354	295.400000	2.863283
1	2.443009	140.471548	3.533324
2	17.067852	63.700000	22.100967
3	4.337363	211.203341	5.190045
4	1.856660	136.500000	2.863283
5	9.932014	189.000000	10.983375
6	3.745587	186.606696	4.589632
7	2.527742	271.600000	3.646212
8	4.518619	63.000000	5.883152
9	4.943637	195.900000	6.576604

	rain_accumulation_9am	rain_duration_9am	relative_humidity_9am	\
0	0.00	0.0	42.420000	
1	0.00	0.0	24.328697	
2	0.00	20.0	8.900000	
3	0.00	0.0	12.189102	
4	8.90	14730.0	92.410000	
5	0.02	170.0	35.130000	
6	0.00	0.0	10.657422	
7	0.00	0.0	80.470000	
8	0.00	0.0	29.580000	
9	0.00	0.0	88.600000	

	relative_humidity_3pm
0	36.160000
1	19.426597
2	14.460000
3	12.742547
4	76.740000
5	33.930000
6	21.385657
7	74.920000
8	24.030000
9	68.050000

Comenzamos a revisar la información y a limpiar información inservible, también revisamos que variables nos interesan

```
[8]: data[data.isnull().any(axis=1)]
```

```
[8]:
```

	number	air_pressure_9am	air_temp_9am	avg_wind_direction_9am	\
16	16	917.890000	NaN	169.200000	
111	111	915.290000	58.820000	182.600000	
177	177	915.900000	NaN	183.300000	
262	262	923.596607	58.380598	47.737753	
277	277	920.480000	62.600000	194.400000	
334	334	916.230000	75.740000	149.100000	
358	358	917.440000	58.514000	55.100000	

361	361	920.444946	65.801845	49.823346
381	381	918.480000	66.542000	90.900000
409	409	NaN	67.853833	65.880616
517	517	920.570000	53.600000	100.100000
519	519	916.250000	55.670000	176.400000
546	546	NaN	42.746000	251.100000
620	620	921.200000	56.786000	192.300000
625	625	912.400000	50.774000	171.600000
656	656	920.830000	66.344000	NaN
670	670	910.920000	48.362000	156.500000
672	672	922.448945	72.863773	NaN
705	705	911.900000	59.072000	199.800000
731	731	922.970166	51.391847	33.810942
737	737	917.895130	76.804690	104.771020
788	788	917.923442	73.249717	42.101739
840	840	918.043767	NaN	181.774042
848	848	915.250000	37.562000	246.500000
861	861	919.065408	NaN	172.303728
869	869	NaN	45.104000	259.000000
998	998	914.140000	71.240000	NaN
1031	1031	922.669195	NaN	47.946284
1035	1035	919.670000	77.576000	171.800000
1063	1063	917.300185	65.790001	NaN
1066	1066	919.564869	73.726732	68.704694

	avg_wind_speed_9am	max_wind_direction_9am	max_wind_speed_9am	\
16	2.192201	196.800000	2.930391	
111	15.613841	189.000000	NaN	
177	4.719943	189.900000	5.346287	
262	10.636273	67.145843	13.671423	
277	2.751436	NaN	3.869906	
334	2.751436	187.500000	4.183078	
358	10.021491	NaN	12.705819	
361	21.520177	61.886944	25.549112	
381	3.467257	89.400000	4.406772	
409	4.328594	78.570923	5.216734	
517	4.697574	NaN	6.285801	
519	6.666081	188.200000	NaN	
546	12.929513	274.400000	17.604718	
620	9.551734	201.400000	11.005745	
625	NaN	181.400000	4.831790	
656	15.457255	189.400000	16.486248	
670	NaN	177.500000	16.128337	
672	3.682370	214.196160	4.849450	
705	1.275056	239.500000	1.834291	
731	NaN	59.290089	11.111555	
737	1.632705	97.178763	NaN	

788	4.132698	64.284969	5.345258
840	0.964376	185.618601	1.570007
848	11.587349	258.700000	NaN
861	2.639600	193.058141	3.326949
869	3.265932	275.000000	4.026492
998	1.722444	232.900000	2.326418
1031	7.969686	65.770066	10.262337
1035	6.554234	191.000000	8.164831
1063	1.879553	222.498226	2.692862
1066	3.551777	102.571616	4.861315

	rain_accumulation_9am	rain_duration_9am	relative_humidity_9am	\
16	0.000	0.000000	48.990000	
111	0.000	0.000000	21.500000	
177	0.000	0.000000	29.260000	
262	0.000	NaN	17.990876	
277	0.000	0.000000	52.580000	
334	NaN	1480.000000	31.880000	
358	0.000	0.000000	13.880000	
361	NaN	40.364018	12.278715	
381	NaN	0.000000	20.640000	
409	0.000	0.000000	18.487385	
517	4.712	14842.000000	79.880000	
519	0.000	0.000000	72.550000	
546	14.627	7825.000000	87.870000	
620	NaN	0.000000	59.790000	
625	0.000	0.000000	86.840000	
656	0.000	0.000000	23.770000	
670	4.970	10560.000000	80.560000	
672	0.000	0.000000	16.753670	
705	NaN	0.000000	77.630000	
731	0.000	4.735034	34.807753	
737	0.000	0.000000	13.771311	
788	0.000	NaN	6.939692	
840	0.000	0.000000	11.911222	
848	3.171	2891.000000	91.000000	
861	0.000	0.000000	12.497839	
869	0.000	80.000000	85.270000	
998	0.000	0.000000	24.200000	
1031	0.000	0.000000	18.920805	
1035	0.000	NaN	56.860000	
1063	0.000	0.000000	14.972668	
1066	NaN	0.000000	11.657314	

	relative_humidity_3pm
16	51.190000
111	29.690000

177	46.500000
262	16.461685
277	54.030000
334	32.900000
358	25.930000
361	7.618649
381	14.350000
409	20.356594
517	84.530000
519	74.390000
546	70.770000
620	77.750000
625	64.740000
656	51.630000
670	88.220000
672	17.804720
705	59.130000
731	18.418179
737	16.792455
788	18.793825
840	18.154358
848	90.780000
861	13.438518
869	90.260000
998	41.380000
1031	19.641841
1035	50.650000
1063	20.966267
1066	17.331823

Comenzamos borrando los valores nulos y eliminamos la columna number que no nos serviría para el ejercicio...

```
[9]: data = data.dropna()
del data['number']
```

```
[10]: data.head(20)
```

```
[10]:
```

	air_pressure_9am	air_temp_9am	avg_wind_direction_9am	\
0	918.060000	74.822000	271.100000	
1	917.347688	71.403843	101.935179	
2	923.040000	60.638000	51.000000	
3	920.502751	70.138895	198.832133	
4	921.160000	44.294000	277.800000	
5	915.300000	78.404000	182.800000	
6	915.598868	70.043304	177.875407	
7	918.070000	51.710000	242.400000	
8	920.080000	80.582000	40.700000	

9	915.010000	47.498000	163.100000
10	919.650000	77.036000	70.600000
11	915.640000	45.716000	241.600000
12	917.390000	49.784000	204.100000
13	920.820000	62.438000	213.600000
14	911.000000	86.432000	202.900000
15	922.383131	70.865263	36.174175
17	916.915255	77.018961	234.539345
18	918.800000	67.082000	176.100000
19	922.040000	68.576000	58.300000
20	919.992262	62.964383	54.799094

	avg_wind_speed_9am	max_wind_direction_9am	max_wind_speed_9am \
0	2.080354	295.400000	2.863283
1	2.443009	140.471548	3.533324
2	17.067852	63.700000	22.100967
3	4.337363	211.203341	5.190045
4	1.856660	136.500000	2.863283
5	9.932014	189.000000	10.983375
6	3.745587	186.606696	4.589632
7	2.527742	271.600000	3.646212
8	4.518619	63.000000	5.883152
9	4.943637	195.900000	6.576604
10	3.825167	85.500000	4.764682
11	5.860783	265.800000	8.030615
12	1.275056	211.800000	2.013246
13	2.617220	165.700000	3.310671
14	1.207948	162.900000	1.677705
15	1.847278	58.428632	2.529142
17	2.274725	229.474199	2.906513
18	4.876529	183.400000	5.569981
19	9.551734	81.900000	12.571603
20	12.680436	74.254223	15.452306

	rain_accumulation_9am	rain_duration_9am	relative_humidity_9am \
0	0.00	0.0	42.420000
1	0.00	0.0	24.328697
2	0.00	20.0	8.900000
3	0.00	0.0	12.189102
4	8.90	14730.0	92.410000
5	0.02	170.0	35.130000
6	0.00	0.0	10.657422
7	0.00	0.0	80.470000
8	0.00	0.0	29.580000
9	0.00	0.0	88.600000
10	0.00	0.0	22.070000
11	0.55	1770.0	90.560000

12	0.00	0.0	73.150000
13	0.00	0.0	43.640000
14	0.00	0.0	15.190000
15	0.00	0.0	12.110889
17	0.00	0.0	21.031462
18	0.00	0.0	18.900000
19	0.00	0.0	7.540000
20	0.00	0.0	18.809518

	relative_humidity_3pm
0	36.160000
1	19.426597
2	14.460000
3	12.742547
4	76.740000
5	33.930000
6	21.385657
7	74.920000
8	24.030000
9	68.050000
10	32.130000
11	79.090000
12	58.430000
13	27.990000
14	24.370000
15	14.801706
17	20.755683
18	45.870000
19	7.740000
20	14.649909

Ahora vamos a seleccionar cuales columnas van a ser nuestros features y cuales serían nuestro target, en este caso nuestro target sería `relative_humidity_3pm` y el resto de columnas serían nuestros features, por otro lado, haciendo un análisis visual, vemos que las columnas `rain_acumulation_9am` y `rain_duration_9am` tiene casi todos los valores en 0 lo que hace pensar que durante el tiempo de estudio no hubo mucha lluvia y estas características no están influyendo de manera relevante en el cálculo...

```
[11]: features = [
    → ['air_temp_9am', 'avg_wind_direction_9am', 'avg_wind_speed_9am', 'max_wind_direction_9am', 'max_
target = ['relative_humidity_3pm']
x = data[features]
y = data[target]
```

Vemos como quedan las features...

```
[12]: x.head(10)
```

```
[12]:  air_temp_9am  avg_wind_direction_9am  avg_wind_speed_9am  \
0      74.822000      271.100000      2.080354
1      71.403843      101.935179      2.443009
2      60.638000      51.000000      17.067852
3      70.138895      198.832133      4.337363
4      44.294000      277.800000      1.856660
5      78.404000      182.800000      9.932014
6      70.043304      177.875407      3.745587
7      51.710000      242.400000      2.527742
8      80.582000      40.700000      4.518619
9      47.498000      163.100000      4.943637

      max_wind_direction_9am  max_wind_speed_9am  relative_humidity_9am
0          295.400000      2.863283      42.420000
1          140.471548      3.533324      24.328697
2           63.700000     22.100967      8.900000
3          211.203341      5.190045     12.189102
4          136.500000      2.863283     92.410000
5          189.000000     10.983375     35.130000
6          186.606696      4.589632     10.657422
7          271.600000      3.646212     80.470000
8           63.000000      5.883152     29.580000
9          195.900000      6.576604     88.600000
```

Ahora el target...

```
[13]: y.head(10)
```

```
[13]:  relative_humidity_3pm
0          36.160000
1          19.426597
2          14.460000
3          12.742547
4          76.740000
5          33.930000
6          21.385657
7          74.920000
8          24.030000
9          68.050000
```


Se realiza una revisión visual de las features respecto al target con el fin de verificar si se ve alguna correlación entre ellas. Primero con la presión del aire a las 9am (este se elimina después de hacer pruebas ya que elevaba bastante el error cuadrático medio)...

```
[14]: #plt.scatter(x['air_pressure_9am'],y,color="darkgreen", label="Data", alpha=.1)
```

Ahora con la temperatura del aire a las 9am...


```
[15]: plt.scatter(x['air_temp_9am'],y,color="darkgreen", label="Data", alpha=.1)
```


[15]: <matplotlib.collections.PathCollection at 0x227c93cb460>




```
[16]: plt.scatter(x['avg_wind_direction_9am'],y,color="darkgreen", label="Data",  
→alpha=.1)
```

[16]: <matplotlib.collections.PathCollection at 0x227c9466eb0>




```
[17]: plt.scatter(x['avg_wind_speed_9am'],y,color="darkgreen", label="Data", alpha=.1)
```

```
[17]: <matplotlib.collections.PathCollection at 0x227c94c5ee0>
```




```
[18]: plt.scatter(x['max_wind_direction_9am'],y,color="darkgreen", label="Data",  
↪alpha=.1)
```

```
[18]: <matplotlib.collections.PathCollection at 0x227c9517880>
```




```
[19]: plt.scatter(x['max_wind_speed_9am'],y,color="darkgreen", label="Data", alpha=.1)
```

```
[19]: <matplotlib.collections.PathCollection at 0x227c9567580>
```



```
[20]: plt.scatter(x['relative_humidity_9am'],y,color="darkgreen", label="Data",  
↳alpha=.1)
```

```
[20]: <matplotlib.collections.PathCollection at 0x227c95b6760>
```



Separamos los datos de entrenamiento y de prueba, manejando una relación de 70% de entrenamiento y 30% para testing

```
[21]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30,  
↳random_state=324)
```

Debido a que el rango de valores entre las features es un poco diferente, podemos utilizar un escalador para homogeneizar los datos

```
[22]: scale = preprocessing.StandardScaler()  
scale.fit(x_train)  
x_train = scale.transform(x_train)
```

Comenzamos con el entrenamiento del modelo, al cual le pasamos los datos de entrenamiento...

```
[23]: regressor = LinearRegression()  
regressor.fit(x_train, y_train)
```

```
[23]: LinearRegression()
```

Ahora vamos al testing y le pasamos los datos de testeo...

```
[24]: x_test=scale.transform(x_test)
      y_prediction = regressor.predict(x_test)
```

Calculamos en error cuadrático medio para ver que tan precisa fue la predicción...

```
[25]: RMSE = sqrt(mean_squared_error(y_true=y_test, y_pred = y_prediction))
      regressor.score(x_test, y_test)
```

```
[25]: 0.814604054492224
```

```
[26]: print ( RMSE )
```

```
9.569374524572197
```


```
[27]: y_test.shape
```

```
[27]: (320, 1)
```

```
[28]: y_prediction.shape
```

```
[28]: (320, 1)
```

```
[29]: n = len(y_test)
      t = np.array(range(n))
      plt.figure(figsize=(15, 4))
      plt.plot(t, y_test, label="Target de prueba")
      plt.plot(t, y_prediction, label="Target predecido")
      plt.legend()
      plt.xlabel("n")
      plt.ylabel("Target")
      plt.show()
```



```
[ ]:
```